


Preliminary Questions

Xarxes, Grau d'Enginyeria Informàtica

Document version: 1.0
Date: Feb. 2020



UAB

Universitat Autònoma
de Barcelona



Departament d'Enginyeria
de la Informació
i de les Comunicacions

Document Information

Document identifier: preliminary_questions.md
Title: Preliminary Questions
Number of pages: 5
Date: Feb. 2020
Author(s): Ian Blanes, Rubén Martínez-Vidal, Angel Elbaz, Kevin Chow.

Change Control

Revision	Date	Changes
1.0	Feb. 2020	(initial version)

Contents

1	Introduction	2
2	How to answer the questions	2
3	Deliverable	2
4	Questions	2
4.1	Data types and language issues	2
4.2	Sockets	2
4.3	Error handling	3
4.4	Binding	3
4.5	Memory problems	3
4.6	Listening and connecting	3
4.7	Fork server	3
4.8	Non-blocking input and output	4
4.9	Polling	4
4.10	Advanced	4
5	Examples	4
5.1	Example 1	4
5.1.1	Incorrect answer	4
5.1.2	Incorrect answer	4
5.1.3	Correct answer	5
5.2	Example 2	5
5.2.1	Incorrect answer	5
5.2.2	Correct answer	5
5.3	Example 3	5
5.3.1	Correct answer	5

1 Introduction

This document contains various questions related to the **POSIX socket API**. Students should carefully study and answer these questions. After successfully completing this assignment, students should be knowledgeable of most of the quirks and details of the API.

2 How to answer the questions

These questions must be answered having application *portability* in mind. This means that the correct answer is the one that does not rely on the details of any specific implementation.

Most answers SHOULD only contain information from authoritative sources. The authoritative source will most likely be the POSIX.1-2017 specification available at <https://pubs.opengroup.org/onlinepubs/9699919799/>.

Tutorials and blogs often describe specific implementations and often contain small mistakes. They may be useful to help figure out how to answer some questions. However, these sources are not authoritative, and the equivalent information needs to be sought in authoritative sources.

However, there will be cases where authoritative sources are not available. Such as questions 4 and 5.

A few examples are available at the end of this document.

3 Deliverable

The deliverable for this assignment is **a single .odt file** to be delivered in the Campus Virtual.

NOTES:

- A template is provided.
- It is not necessary that all members of a group contribute to the writing of each answer.
- **Each answer must be verified** to ensure its correctness by some group member, preferably, someone different than the one who wrote it.
- **Each answer must include information of the person who wrote it and the (other) person who verified it** (see template).

4 Questions

4.1 Data types and language issues

1. What is the function `htons` used for?
2. What is `uint32_t`?
3. What is the difference between `ssize_t` and `size_t`?
4. What is the value of `b` in the following piece of code?

```
const ssize_t a = -1;
const size_t b = (size_t) a;
```

5. What is the value of `b` in the following piece of code?

```
const uint32_t a = 256;
const uint8_t b = (uint8_t) a;
```

4.2 Sockets

6. What is a *file descriptor*?
7. What values for the `protocol` parameter of the `socket` function are available for `AF_INET` in POSIX? What does it mean that the default value for `protocol` in `socket` is *implementation defined*?
8. Give a piece of code that creates a TCP socket.

9. Give a piece of code that creates a UDP socket.
10. How do we close a socket?
11. What happens if we leave it open?
12. If we never close any sockets, we may eventually receive an EMFILE error code in socket. Why?

4.3 Error handling

13. What is errno? How is it used?
14. Why is the perror function used? Show a small example.
15. Can we modify errno ourselves in our code? Should we?
16. Which errors can result from a socket function call? And from sendto? (do not describe them the error codes, just mention them: EAGAIN)
17. What happens if a program terminates (e.g., exit(EXIT_FAILURE)) while still having some open sockets (not having called close)?

4.4 Binding

18. What does bind do?
19. What is the difference between sockaddr, sockaddr_in, sockaddr_in6, and {sockaddr_storage}?
20. Why is there an addr_len parameter in bind?
21. What is in_port_t?
22. Which of the following pieces of code regarding port 8080 is the right one: in_port_t port = htons(8080), in_port_t port = ntohs(8080), or in_port_t port = 8080? Why?
23. Give a proper initialization code for a sockaddr_in structure for port 8080 and address 127.0.0.1.
24. Give a proper initialization code for a sockaddr_in structure for port 8080 and any address.

4.5 Memory problems

25. What is a null pointer?
26. What does it mean when a program receives a SIGSEGV?
27. What is valgrind and how can we use it to debug memory problems?
28. Why is it a good practice to free all allocated memory before calling exit(EXIT_SUCCESS)? (hint: valgrind)

4.6 Listening and connecting

29. What does connect do?
30. Why does connect take one addr parameter?
31. Why do we need to call accept after listen?
32. Why does accept take one addr parameter?
33. In the following piece of code, are a and b equal? Why?

```
int b = accept(a, ...);
```

34. What happens if we listen or connect to an unbound socket? (bind has not been called).

4.7 Fork server

35. On a forking server, after accept and fork, why should the parent close the socket returned by accept? (hint: resource exhaustion)

4.8 Non-blocking input and output

36. What is non-blocking I/O?
37. How can we set a socket to perform non-blocking I/O? (give code)
38. What happens with a read on a blocking socket if all the requested data is not available?
39. What happens with a write on a blocking socket if the socket buffer is full?
40. What happens with a read on a non-blocking socket if all the requested data is not available? (hint: there are two possibilities)
41. What happens with a write on a non-blocking socket if the data to write does not fit in the socket buffer? (hint: there are two possibilities)

4.9 Polling

42. How is the function `poll` used?
43. Can we mix file descriptors for sockets and disk files in a single `poll` call?
44. Which events can report a `poll` call for `fds[i].events = POLLIN`?

4.10 Advanced

45. Are socket options inherited from a listening socket with `accept`?
46. Which size is the option for `SO_RCVLOWAT`?
47. After `poll` signals `revents = POLLOUT` on a non-blocking socket, will a very large send block execution?
48. Can `getaddrinfo` return an empty linked list?

5 Examples

5.1 Example 1

Question: Which header files are necessary for the socket function?

5.1.1 Incorrect answer

The following header files are required:

```
#include <sys/socket.h>
#include <netinet/in.h>
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <arpa/inet.h>
//#include <fcntl.h>
#include <unistd.h>
```

Source: <https://stackoverflow.com/questions/9198608/simple-socket-programming-code-working>

5.1.2 Incorrect answer

Headers `sys/types.h` and `sys/socket.h` are needed.

Source: <http://man7.org/linux/man-pages/man2/socket.2.html>

5.1.3 Correct answer

Only the inclusion of the `sys/socket.h` headers is necessary. However, the `sys/types.h` header was required for some historical implementations.

Sources: https://pubs.opengroup.org/onlinepubs/9699919799/basedefs/sys_socket.h.html and <http://man7.org/linux/man-pages/man2/socket.2.html>

5.2 Example 2

Question: Which value shall we use for the domain field while creating an IPv4 socket? `PF_INET` or `AF_INET`?

5.2.1 Incorrect answer

We should use `PF_INET` in the `socket()` call because it refers to a protocol, and `AF_INET` in the `addrsock_in` structure because it refers to an address.

Source: <https://stackoverflow.com/questions/6729366/what-is-the-difference-between-af-inet-and-pf-inet-in-socket-programming>

5.2.2 Correct answer

The constant `PF_INET` is not defined in POSIX, and we should not use it. The constant `AF_INET` is defined as

Internet domain sockets for use with IPv4 addresses.

We must use `AF_INET`.

Sources: <https://pubs.opengroup.org/onlinepubs/009695399/basedefs/sys/socket.h.html>

5.3 Example 3

Question: Why is the following piece of code incorrect?

```
//...
static const uint8_t MSG_OK = 1;
uint8_t buffer[8];
recv(sockfd1, buffer, 8, 0);

if (buffer[0] == MSG_OK) {
    // Do something
} else {
    // Do something else
}
```

5.3.1 Correct answer

The function `recv` returns an `ssize_t` value that must be checked for errors before assuming that the function succeeded and that `buffer[0]` contains valid data.

Source: <https://pubs.opengroup.org/onlinepubs/9699919799.2016edition/functions/recv.html>